

16 — heads-averaged — Cumulative attention rollout matrix per layer (50×50: 1 CLS + 7×7 patches)

